

Apache HTTP Server Version 2.4

Apache Module mod_proxy

Description:	Multi-protocol proxy/gateway server
Status:	Extension
Module Identifier:	proxy_module
Source File:	mod_proxy.c

Summary

Warning
Do not enable proxying with **ProxyRequests** until you have secured your server ([↗ #access](#)) . Open proxy servers are dangerous both to your network and to the Internet at large.

mod_proxy and related modules implement a proxy/gateway for Apache HTTP Server, supporting a number of popular protocols as well as several different load balancing algorithms. Third-party modules can add support for additional protocols and load balancing algorithms.

A set of modules must be loaded into the server to provide the necessary features. These modules can be included statically at build time or dynamically via the **LoadModule** directive). The set must include:

- mod_proxy**, which provides basic proxy capabilities
- mod_proxy_balancer** and one or more balancer modules if load balancing is required. (See **mod_proxy_balancer** for more information.)
- one or more proxy scheme, or protocol, modules:

Protocol	Module
AJP13 (Apache JServe Protocol version 1.3)	mod_proxy_ajp
CONNECT (for SSL)	mod_proxy_connect
FastCGI	mod_proxy_fcgi
ftp	mod_proxy_ftp
HTTP/0.9, HTTP/1.0, and HTTP/1.1	mod_proxy_http
HTTP/2.0	mod_proxy_http2
SCGI	mod_proxy_scgi
UWSGI	mod_proxy_uwsgi
WS and WSS (Web-sockets)	mod_proxy_wstunnel

In addition, extended features are provided by other modules. Caching is provided by **mod_cache** and related modules. The ability to contact remote servers using the SSL/TLS protocol is provided by the SSLProxy* directives of **mod_ssl**. These additional modules will need to be loaded and configured to take advantage of these features.

Forward Proxies and Reverse Proxies/Gateways

Apache HTTP Server can be configured in both a *forward* and *reverse* proxy (also known as *gateway*) mode.

An ordinary *forward proxy* is an intermediate server that sits between the client and the *origin server*. In order to get content from the origin server, the client sends a request to the proxy naming the origin server as the target. The proxy then requests the content from the origin server and returns it to the client. The client must be specially configured to use the forward proxy to access other sites.

A typical usage of a forward proxy is to provide Internet access to internal clients that are otherwise restricted by a firewall. The forward proxy can also use caching (as provided by **mod_cache**) to reduce network usage.

The forward proxy is activated using the **ProxyRequests** directive. Because forward proxies allow clients to access arbitrary sites through your server and to hide their true origin, it is essential that you secure your server ([↗ #access](#)) so that only authorized clients can access the proxy before activating a forward proxy.

A *reverse proxy* (or *gateway*), by contrast, appears to the client just like an ordinary web server. No special configuration on the client is necessary. The client makes ordinary requests for content in the namespace of the reverse proxy. The reverse proxy then decides where to send those requests and returns the content as if it were itself the origin.

A typical usage of a reverse proxy is to provide Internet users access to a server that is behind a firewall. Reverse proxies can also be used to balance load among several back-end servers or to provide caching for a slower back-end server. In addition, reverse proxies can be used simply to bring several servers into the same URL space.

A reverse proxy is activated using the **ProxyPass** directive or the [P] flag to the **RewriteRule** directive. It is **not** necessary to turn **ProxyRequests** on in order to configure a reverse proxy.

Basic Examples

The examples below are only a very basic idea to help you get started. Please read the documentation on the individual directives.

In addition, if you wish to have caching enabled, consult the documentation from **mod_cache**.

Reverse Proxy

```
ProxyPass "/foo" "http://foo.example.com/bar"
ProxyPassReverse "/foo" "http://foo.example.com/bar"
```

Forward Proxy

```
ProxyRequests On
ProxyVia On

<Proxy "*">
  Require host internal.example.com
</Proxy>
```

Websocket Upgrade (2.4.47 and later)

```
ProxyPass "/some/ws/capable/path/" "http://example.com/some/ws/capable/path/" upgrade=websocket
```

Access via Handler

You can also force a request to be handled as a reverse-proxy request, by creating a suitable Handler pass-through. The example configuration below will pass all requests for PHP scripts to the specified FastCGI server using reverse proxy:

Reverse Proxy PHP scripts

```
<FilesMatch "\.php$">
    # Unix sockets require 2.4.7 or later
    SetHandler "proxy:unix:/path/to/app.sock|fcgi://localhost/"
</FilesMatch>
```

This feature is available in Apache HTTP Server 2.4.10 and later.

Workers

The proxy manages the configuration of origin servers and their communication parameters in objects called *workers*. There are two built-in workers: the default forward proxy worker and the default reverse proxy worker. Additional workers can be configured explicitly.

The two default workers have a fixed configuration and will be used if no other worker matches the request. They do not use HTTP Keep-Alive or connection reuse. The TCP connections to the origin server will instead be opened and closed for each request.

Explicitly configured workers are identified by their URL. They are usually created and configured using **ProxyPass** or **ProxyPassMatch** when used for a reverse proxy:

```
ProxyPass "/example" "http://backend.example.com" connectiontimeout=5 timeout=30
```

This will create a worker associated with the origin server URL `http://backend.example.com` that will use the given timeout values. When used in a forward proxy, workers are usually defined via the **ProxySet** directive:

```
ProxySet "http://backend.example.com" connectiontimeout=5 timeout=30
```

or alternatively using **Proxy** and **ProxySet**:

```
<Proxy "http://backend.example.com">
    ProxySet connectiontimeout=5 timeout=30
</Proxy>
```

Using explicitly configured workers in the forward mode is not very common, because forward proxies usually communicate with many different origin servers. Creating explicit workers for some of the origin servers can still be useful if they are used very often. Explicitly configured workers have no concept of forward or reverse proxying by themselves. They encapsulate a common concept of communication with origin servers. A worker created by **ProxyPass** for use in a reverse proxy will also be used for forward proxy requests whenever the URL to the origin server matches the worker URL, and vice versa.

The URL identifying a direct worker is the URL of its origin server including any path components given:

```
ProxyPass "/examples" "http://backend.example.com/examples"
ProxyPass "/docs" "http://backend.example.com/docs"
```

This example defines two different workers, each using a separate connection pool and configuration.

Worker Sharing

Worker sharing happens if the worker URLs overlap, which occurs when the URL of some worker is a leading substring of the URL of another worker defined later in the configuration file. In the following example

```
ProxyPass "/apps" "http://backend.example.com/" timeout=60
ProxyPass "/examples" "http://backend.example.com/examples" timeout=10
```

the second worker isn't actually created. Instead the first worker is used. The benefit is, that there is only one connection pool, so connections are more often reused. Note that all configuration attributes given explicitly for the later worker will be ignored. This will be logged as a warning. In the above example, the resulting timeout value for the URL `/examples` will be 60 instead of 10!

If you want to avoid worker sharing, sort your worker definitions by URL length, starting with the longest worker URLs. If you want to maximize worker sharing, use the reverse sort order. See also the related warning about ordering **ProxyPass** directives.

Explicitly configured workers come in two flavors: *direct workers* and *(load) balancer workers*. They support many important configuration attributes which are described below in the **ProxyPass** directive. The same attributes can also be set using **ProxySet**.

The set of options available for a direct worker depends on the protocol which is specified in the origin server URL. Available protocols include `ajp`, `fcgi`, `ftp`, `http` and `scgi`.

Balancer workers are virtual workers that use direct workers known as their members to actually handle the requests. Each balancer can have multiple members. When it handles a request, it chooses a member based on the configured load balancing algorithm.

A balancer worker is created if its worker URL uses `balancer` as the protocol scheme. The balancer URL uniquely identifies the balancer worker. Members are added to a balancer using **BalancerMember**.

DNS resolution for origin domains

DNS resolution happens when the socket to the origin domain is created for the first time. When connection reuse is enabled, each backend domain is resolved only once per child process, and cached for all further connections until the child is recycled. This information should to be considered while planning DNS maintenance tasks involving backend domains. Please also check **ProxyPass** parameters for more details about connection reuse.

Controlling Access to Your Proxy

You can control who can access your proxy via the `<Proxy>` control block as in the following example:

```
<Proxy "*">
    Require ip 192.168.0
</Proxy>
```

For more information on access control directives, see `mod_authz_host`.

Strictly limiting access is essential if you are using a forward proxy (using the `ProxyRequests` directive). Otherwise, your server can be used by any client to access arbitrary hosts while hiding his or her true identity. This is dangerous both for your network and for the Internet at large. When using a reverse proxy (using the `ProxyPass` directive with `ProxyRequests Off`), access control is less critical because clients can only contact the hosts that you have specifically configured.

See Also the Proxy-Chain-Auth ([↗ mod_proxy_http.html#env](#)) environment variable.

Slow Startup

If you're using the `ProxyBlock` directive, hostnames' IP addresses are looked up and cached during startup for later match test. This may take a few seconds (or more) depending on the speed with which the hostname lookups occur.

Intranet Proxy

An Apache httpd proxy server situated in an intranet needs to forward external requests through the company's firewall (for this, configure the `ProxyRemote` directive to forward the respective *scheme* to the firewall proxy). However, when it has to access resources within the intranet, it can bypass the firewall when accessing hosts. The `NoProxy` directive is useful for specifying which hosts belong to the intranet and should be accessed directly.

Users within an intranet tend to omit the local domain name from their WWW requests, thus requesting "http://somehost/" instead of `http://somehost.example.com/`. Some commercial proxy servers let them get away with this and simply serve the request, implying a configured local domain. When the `ProxyDomain` directive is used and the server is configured for proxy service ([↗ #proxyrequests](#)) , Apache httpd can return a redirect response and send the client to the correct, fully qualified, server address. This is the preferred method since the user's bookmark files will then contain fully qualified hosts.

Protocol Adjustments

For circumstances where `mod_proxy` is sending requests to an origin server that doesn't properly implement keepalives or HTTP/1.1, there are two environment variables ([↗ ../env.html](#)) that can force the request to use HTTP/1.0 with no keepalive. These are set via the `SetEnv` directive.

These are the `force-proxy-request-1.0` and `proxy-nokeepalive` notes.

```
<Location "/buggyappserver/">
    ProxyPass "http://buggyappserver:7001/foo/"
    SetEnv force-proxy-request-1.0 1
    SetEnv proxy-nokeepalive 1
</Location>
```

In 2.4.26 and later, the "no-proxy" environment variable can be set to disable `mod_proxy` processing the current request. This variable should be set with `SetEnvIf`, as `SetEnv` is not evaluated early enough.

Request Bodies

Some request methods such as POST include a request body. The HTTP protocol requires that requests which include a body either use chunked transfer encoding or send a `Content-Length` request header. When passing these requests on to the origin server, `mod_proxy_http` will always attempt to send the `Content-Length`. But if the body is large and the original request used chunked encoding, then chunked encoding may also be used in the upstream request. You can control this selection using environment variables ([↗ ../env.html](#)) . Setting `proxy-sendcl` ensures maximum compatibility with upstream servers by always sending the `Content-Length`, while setting `proxy-sendchunked` minimizes resource usage by using chunked encoding.

Under some circumstances, the server must spool request bodies to disk to satisfy the requested handling of request bodies. For example, this spooling will occur if the original body was sent with chunked encoding (and is large), but the administrator has asked for backend requests to be sent with `Content-Length` or as HTTP/1.0. This spooling can also occur if the request body already has a `Content-Length` header, but the server is configured to filter incoming request bodies.

Reverse Proxy Request Headers

When acting in a reverse-proxy mode (using the `ProxyPass` directive, for example), `mod_proxy_http` adds several request headers in order to pass information to the origin server. These headers are:

X-Forwarded-For
The IP address of the client.

X-Forwarded-Host
The original host requested by the client in the `Host` HTTP request header.

X-Forwarded-Server
The hostname of the proxy server.

Be careful when using these headers on the origin server, since they will contain more than one (comma-separated) value if the original request already contained one of these headers. For example, you can use `%{X-Forwarded-For}i` in the log format string of the origin server to log the original clients IP address, but you may get more than one address if the request passes through several proxies.

See also the `ProxyPreserveHost` and `ProxyVia` directives, which control other request headers.

Note: If you need to specify custom request headers to be added to the forwarded request, use the `RequestHeader` directive.

BalancerGrowth Directive

Description:	Number of additional Balancers that can be added Post-configuration
Syntax:	<code>BalancerGrowth #</code>
Default:	<code>BalancerGrowth 5</code>
Context:	server config, virtual host
Status:	Extension
Module:	<code>mod_proxy</code>
Compatibility:	<code>BalancerGrowth</code> is only available in Apache HTTP Server 2.3.13 and later.

This directive allows for growth potential in the number of Balancers available for a virtualhost in addition to the number pre-configured. It only takes effect if there is at least one pre-configured Balancer.

BalancerInherit Directive

Description:	Inherit ProxyPassed Balancers/Workers from the main server
Syntax:	BalancerInherit On Off
Default:	BalancerInherit On
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy
Compatibility:	BalancerInherit is only available in Apache HTTP Server 2.4.5 and later.

This directive will cause the current server/vhost to "inherit" ProxyPass Balancers and Workers defined in the main server. This can cause issues and inconsistent behavior if using the Balancer Manager and so should be disabled if using that feature.

The setting in the global server defines the default for all vhosts.

BalancerMember Directive

Description:	Add a member to a load balancing group
Syntax:	BalancerMember [<i>balancerurl</i>] <i>url</i> [<i>key=value</i> [<i>key=value ...</i>]]
Context:	directory
Status:	Extension
Module:	mod_proxy
Compatibility:	BalancerMember is only available in Apache HTTP Server 2.2 and later.

This directive adds a member to a load balancing group. It can be used within a <Proxy *balancer:/. . .*> container directive and can take any of the key value pair parameters available to **ProxyPass** directives.

One additional parameter is available only to **BalancerMember** directives: *loadfactor*. This is the member load factor - a decimal number between 1.0 (default) and 100.0, which defines the weighted load to be applied to the member in question.

The *balancerurl* is only needed when not within a <Proxy *balancer:/. . .*> container directive. It corresponds to the url of a balancer defined in **ProxyPass** directive.

The path component of the balancer URL in any <Proxy *balancer:/. . .*> container directive is ignored.

Trailing slashes should typically be removed from the URL of a **BalancerMember**.

BalancerPersist Directive

Description:	Attempt to persist changes made by the Balancer Manager across restarts.
Syntax:	BalancerPersist On Off
Default:	BalancerPersist Off
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy
Compatibility:	BalancerPersist is only available in Apache HTTP Server 2.4.4 and later.

This directive will cause the shared memory storage associated with the balancers and balancer members to be persisted across restarts. This allows these local changes to not be lost during the normal restart/graceful state transitions.

NoProxy Directive

Description:	Hosts, domains, or networks that will be connected to directly
Syntax:	NoProxy <i>host</i> [<i>host</i>] ...
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

This directive is only useful for Apache httpd proxy servers within intranets. The **NoProxy** directive specifies a list of subnets, IP addresses, hosts and/or domains, separated by spaces. A request to a host which matches one or more of these is always served directly, without forwarding to the configured **ProxyRemote** proxy server(s).

Example
ProxyRemote "*" "http://firewall.example.com:81" NoProxy ".example.com" "192.168.112.0/21"

The *host* arguments to the **NoProxy** directive are one of the following type list:

Domain

A *Domain* is a partially qualified DNS domain name, preceded by a period. It represents a list of hosts which logically belong to the same DNS domain or zone (*i.e.*, the suffixes of the hostnames are all ending in *Domain*).

Examples
.com .example.org.

To distinguish *Domains* from *Hostnames* (both syntactically and semantically; a DNS domain can have a DNS A record, too!), *Domains* are always written with a leading period.

Note
Domain name comparisons are done without regard to the case, and <i>Domains</i> are always assumed to be anchored in the root of the DNS tree; therefore, the two domains .EXAmple.com and .example.com. (note the trailing period) are considered equal. Since a domain comparison does not involve a DNS lookup, it is much more efficient than subnet comparison.

SubNet

A *SubNet* is a partially qualified internet address in numeric (dotted quad) form, optionally followed by a slash and the netmask, specified as the number of significant bits in the *SubNet*. It is used to represent a subnet of hosts which can be reached over a common network interface. In the absence of the explicit net mask it is assumed that omitted (or zero valued) trailing digits specify the mask. (In this case, the netmask can only be multiples of 8 bits wide.) Examples:

- 192.168 or 192.168.0.0**
the subnet 192.168.0.0 with an implied netmask of 16 valid bits (sometimes used in the netmask form 255.255.0.0)
- 192.168.112.0/21**
the subnet 192.168.112.0/21 with a netmask of 21 valid bits (also used in the form 255.255.248.0)

As a degenerate case, a *SubNet* with 32 valid bits is the equivalent to an *IPAddr*, while a *SubNet* with zero valid bits (*e.g.*, 0.0.0.0/0) is the same as the constant *_Default_*, matching any IP address.

IPAddr

A *IPAddr* represents a fully qualified internet address in numeric (dotted quad) form. Usually, this address represents a host, but there need not necessarily be a DNS domain name connected with the address.

Example
192.168.123.7

Note
An *IPAddr* does not need to be resolved by the DNS system, so it can result in more effective apache performance.

Hostname

A *Hostname* is a fully qualified DNS domain name which can be resolved to one or more *IPAddrs* via the DNS domain name service. It represents a logical host (in contrast to *Domains*, see above) and must be resolvable to at least one *IPAddr* (or often to a list of hosts with different *IPAddrs*).

Examples
prep.ai.example.edu
www.example.org

Note
In many situations, it is more effective to specify an *IPAddr* in place of a *Hostname* since a DNS lookup can be avoided. Name resolution in Apache httpd can take a remarkable deal of time when the connection to the name server uses a slow PPP link.
Hostname comparisons are done without regard to the case, and *Hostnames* are always assumed to be anchored in the root of the DNS tree; therefore, the two hosts `WWW.ExAmPle.com` and `www.example.com.` (note the trailing period) are considered equal.

See also

- DNS Issues

<Proxy> Directive

Description:	Container for directives applied to proxied resources
Syntax:	<Proxy <i>wildcard-url</i> > ...</Proxy>
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

Directives placed in **<Proxy>** sections apply only to matching proxied content. Shell-style wildcards are allowed.

For example, the following will allow only hosts in `yournetwork.example.com` to access content via your proxy server:

```
<Proxy "*">
  Require host yournetwork.example.com
</Proxy>
```

The following example will process all files in the `foo` directory of `example.com` through the `INCLUDES` filter when they are sent through the proxy server:

```
<Proxy "http://example.com/foo/*">
  SetOutputFilter INCLUDES
</Proxy>
```

Differences from the Location configuration section

A backend URL matches the configuration section if it begins with the the *wildcard-url* string, even if the last path segment in the directive only matches a prefix of the backend URL. For example, `<Proxy "http://example.com/foo">` matches all of `http://example.com/foo`, `http://example.com/foo/bar`, and `http://example.com/foobar`. The matching of the final URL differs from the behavior of the **<Location>** section, which for purposes of this note treats the final path component as if it ended in a slash. For more control over the matching, see **<ProxyMatch>**.

See also

- <ProxyMatch>**

Proxy100Continue Directive

Description:	Forward 100-continue expectation to the origin server
Syntax:	Proxy100Continue Off On
Default:	Proxy100Continue On
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy

Compatibility: Available in version 2.4.40 and later

This directive determines whether the proxy should forward 100-continue *Expect*:ation to the origin server and thus let it decide when/if the HTTP request body should be read, or when `Off` the proxy should generate *100 Continue* intermediate response by itself before forwarding the request body.

Effectiveness

This option is of use only for HTTP proxying, as handled by `mod_proxy_http`.

ProxyAddHeaders Directive

Description:	Add proxy information in X-Forwarded-* headers
Syntax:	<code>ProxyAddHeaders Off On</code>
Default:	<code>ProxyAddHeaders On</code>
Context:	server config, virtual host, directory
Status:	Extension
Module:	<code>mod_proxy</code>
Compatibility:	Available in version 2.3.10 and later

This directive determines whether or not proxy related information should be passed to the backend server through X-Forwarded-For, X-Forwarded-Host and X-Forwarded-Server HTTP headers.

Effectiveness

This option is of use only for HTTP proxying, as handled by `mod_proxy_http`.

ProxyBadHeader Directive

Description:	Determines how to handle bad header lines in a response
Syntax:	<code>ProxyBadHeader IsError Ignore StartBody</code>
Default:	<code>ProxyBadHeader IsError</code>
Context:	server config, virtual host
Status:	Extension
Module:	<code>mod_proxy</code>

The **ProxyBadHeader** directive determines the behavior of `mod_proxy` if it receives syntactically invalid response header lines (*i.e.* containing no colon) from the origin server. The following arguments are possible:

IsError
Abort the request and end up with a 502 (Bad Gateway) response. This is the default behavior.

Ignore
Treat bad header lines as if they weren't sent.

StartBody
When receiving the first bad header line, finish reading the headers and treat the remainder as body. This helps to work around buggy backend servers which forget to insert an empty line between the headers and the body.

ProxyBlock Directive

Description:	Words, hosts, or domains that are banned from being proxied
Syntax:	<code>ProxyBlock * word host domain [word host domain] ...</code>
Context:	server config, virtual host
Status:	Extension
Module:	<code>mod_proxy</code>

The **ProxyBlock** directive specifies a list of words, hosts and/or domains, separated by spaces. HTTP, HTTPS, and FTP document requests to sites whose names contain matched words, hosts or domains are *blocked* by the proxy server. The proxy module will also attempt to determine IP addresses of list items which may be hostnames during startup, and cache them for match test as well. That may slow down the startup time of the server.

Example

```
ProxyBlock "news.example.com" "auctions.example.com" "friends.example.com"
```

Note that `example` would also be sufficient to match any of these sites.

Hosts would also be matched if referenced by IP address.

Note also that

```
ProxyBlock "*" "
```

blocks connections to all sites.

ProxyDomain Directive

Description:	Default domain name for proxied requests
Syntax:	<code>ProxyDomain Domain</code>
Context:	server config, virtual host
Status:	Extension
Module:	<code>mod_proxy</code>

This directive is only useful for Apache httpd proxy servers within intranets. The **ProxyDomain** directive specifies the default domain which the apache proxy server will belong to. If a request to a host without a domain name is encountered, a redirection response to the same host with the configured *Domain* appended will be generated.

Example

```
ProxyRemote    "*"    "http://firewall.example.com:81"
NoProxy       ".example.com" "192.168.112.0/21"
ProxyDomain   ".example.com"
```

ProxyErrorOverride Directive

Description:	Override error pages for proxied content
Syntax:	ProxyErrorOverride Off On [<i>code</i> ...]
Default:	ProxyErrorOverride Off
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy
Compatibility:	The list of status codes was added in 2.4.47

This directive is useful for reverse-proxy setups where you want to have a common look and feel on the error pages seen by the end user. This also allows for included files (via `mod_include`'s SSI) to get the error code and act accordingly. (Default behavior would display the error page of the proxied server. Turning this on shows the SSI Error message.)

This directive does not affect the processing of informational (1xx), normal success (2xx), or redirect (3xx) responses.

By default `ProxyErrorOverride` affects all responses with codes between 400 (including) and 600 (excluding).

Example for default behavior

```
ProxyErrorOverride On
```

To change the default behavior, you can specify the status codes to consider, separated by spaces. If you do so, all other status codes will be ignored. You can only specify status codes, that are considered error codes: between 400 (including) and 600 (excluding).

Example for custom status codes

```
ProxyErrorOverride On 403 405 500 501 502 503 504
```

ProxyIOBufferSize Directive

Description:	Determine size of internal data throughput buffer
Syntax:	ProxyIOBufferSize <i>bytes</i>
Default:	ProxyIOBufferSize 8192
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

The `ProxyIOBufferSize` directive adjusts the size of the internal buffer which is used as a scratchpad for the data between input and output. The size must be at least 512.

In almost every case, there's no reason to change that value.

If used with AJP, this directive sets the maximum AJP packet size in bytes. Values larger than 65536 are set to 65536. If you change it from the default, you must also change the `packetSize` attribute of your AJP connector on the Tomcat side! The attribute `packetSize` is only available in Tomcat 5.5.20+ and 6.0.2+

Normally it is not necessary to change the maximum packet size. Problems with the default value have been reported when sending certificates or certificate chains.

<ProxyMatch> Directive

Description:	Container for directives applied to regular-expression-matched proxied resources
Syntax:	<ProxyMatch <i>regex</i> > ...</ProxyMatch>
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

The `<ProxyMatch>` directive is identical to the `<Proxy>` directive, except that it matches URLs using [regular expressions](#) (↗ [../glossary.html#regex](#)) .

From 2.4.8 onwards, named groups and backreferences are captured and written to the environment with the corresponding name prefixed with "MATCH_" and in upper case. This allows elements of URLs to be referenced from within expressions (↗ [../expr.html](#)) and modules like `mod_rewrite`. In order to prevent confusion, numbered (unnamed) backreferences are ignored. Use named groups instead.

```
<ProxyMatch "^http://(?<sitename>[/]+)">
  Require ldap-group cn=%{env:MATCH_SITENAME},ou=combined,o=Example
</ProxyMatch>
```

See also

- `<Proxy>`

ProxyMaxForwards Directive

Description:	Maximum number of proxies that a request can be forwarded through
Syntax:	ProxyMaxForwards <i>number</i>
Default:	ProxyMaxForwards -1
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy
Compatibility:	Default behaviour changed in 2.2.7

The **ProxyMaxForwards** directive specifies the maximum number of proxies through which a request may pass if there's no **Max-Forwards** header supplied with the request. This may be set to prevent infinite proxy loops or a DoS attack.

Example

ProxyMaxForwards 15

Note that setting **ProxyMaxForwards** is a violation of the HTTP/1.1 protocol (RFC2616), which forbids a Proxy setting **Max-Forwards** if the Client didn't set it. Earlier Apache httpd versions would always set it. A negative **ProxyMaxForwards** value, including the default -1, gives you protocol-compliant behavior but may leave you open to loops.

ProxyPass Directive

Description:	Maps remote servers into the local server URL-space
Syntax:	ProxyPass [<i>path</i>] ! <i>url</i> [<i>key=value</i> [<i>key=value</i> ...]] [<i>nocanon</i>] [<i>interpolate</i>] [<i>noquery</i>]
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy
Compatibility:	Unix Domain Socket (UDS) support added in 2.4.7

This directive allows remote servers to be mapped into the space of the local server. The local server does not act as a proxy in the conventional sense but appears to be a mirror of the remote server. The local server is often called a *reverse proxy* or *gateway*. The *path* is the name of a local virtual path; *url* is a partial URL for the remote server and cannot include a query string.

It is strongly suggested to review the concept of a Worker before proceeding any further with this section.

This directive is not supported within **<Directory>**, **<If>** and **<Files>** containers.

The **ProxyRequests** directive should usually be set **off** when using **ProxyPass**.

In 2.4.7 and later, support for using a Unix Domain Socket is available by using a target which prepends `unix:/path/lis.sock|`. For example, to proxy HTTP and target the UDS at /home/www.socket, you would use `unix:/home/www.socket|http://localhost/whatever/`.

Note: The path associated with the `unix:` URL is **DefaultRuntimeDir** aware.

When used inside a **<Location>** section, the first argument is omitted and the local directory is obtained from the **<Location>**. The same will occur inside a **<LocationMatch>** section; however, ProxyPass does not interpret the regexp as such, so it is necessary to use **ProxyPassMatch** in this situation instead.

Suppose the local server has address `http://example.com/`; then

```
<Location "/mirror/foo/">
    ProxyPass "http://backend.example.com/"
</Location>
```

will cause a local request for `http://example.com/mirror/foo/bar` to be internally converted into a proxy request to `http://backend.example.com/bar`.

If you require a more flexible reverse-proxy configuration, see the **RewriteRule** directive with the `[P]` flag.

The following alternative syntax is possible; however, it can carry a performance penalty when present in very large numbers. The advantage of the below syntax is that it allows for dynamic control via the Balancer Manager ([↗ mod_proxy_balancer.html#balancer_manager](#)) interface:

```
ProxyPass "/mirror/foo/" "http://backend.example.com/"
```

If the first argument ends with a trailing `/`, the second argument should also end with a trailing `/`, and vice versa. Otherwise, the resulting requests to the backend may miss some needed slashes and do not deliver the expected results.

The **!** directive is useful in situations where you don't want to reverse-proxy a subdirectory, *e.g.*

```
<Location "/mirror/foo/">
    ProxyPass "http://backend.example.com/"
</Location>
<Location "/mirror/foo/i">
    ProxyPass "!"
</Location>
```

```
ProxyPass "/mirror/foo/i" "!"
ProxyPass "/mirror/foo" "http://backend.example.com"
```

will proxy all requests to `/mirror/foo` to `backend.example.com` *except* requests made to `/mirror/foo/i`.

Mixing ProxyPass settings in different contexts does not work:

```
ProxyPass "/mirror/foo/i" "!"
<Location "/mirror/foo/">
    ProxyPass "http://backend.example.com/"
</Location>
```

In this case, a request to `/mirror/foo/i` will get proxied, because the **ProxyPass** directive in the Location block will be evaluated first. The fact that **ProxyPass** supports both server and directory contexts does not mean that their scope and position in the configuration file will guarantee any ordering or override.

Ordering ProxyPass Directives

The configured **ProxyPass** and **ProxyPassMatch** rules are checked in the order of configuration. The first rule that matches wins. So usually you should sort conflicting **ProxyPass** rules starting with the longest URLs first. Otherwise, later rules for longer URLs will be hidden by any earlier rule which uses a leading substring of the URL.

Note that there is some relation with worker sharing.

Ordering ProxyPass Directives in Locations

Only one **ProxyPass** directive can be placed in a **Location** block, and the most specific location will take precedence.

Exclusions and the no-proxy environment variable

Exclusions must come *before* the general **ProxyPass** directives. In 2.4.26 and later, the "no-proxy" environment variable is an alternative to exclusions, and is the only way to configure an exclusion of a **ProxyPass** directive in **Location** context. This variable should be set with **SetEnvIf**, as **SetEnv** is not evaluated early enough.

ProxyPass key=value Parameters

In Apache HTTP Server 2.1 and later, mod_proxy supports pooled connections to a backend server. Connections created on demand can be retained in a pool for future use. Limits on the pool size and other settings can be coded on the **ProxyPass** directive using **key=value** parameters, described in the tables below.

Maximum connections to the backend

By default, mod_proxy will allow and retain the maximum number of connections that could be used simultaneously by that web server child process. Use the **max** parameter to reduce the number from the default. The pool of connections is maintained per web server child process, and **max** and other settings are not coordinated among all child processes, except when only one child process is allowed by configuration or MPM design.

Use the **ttl** parameter to set an optional time to live; connections which have been unused for at least **ttl** seconds will be closed. **ttl** can be used to avoid using a connection which is subject to closing because of the backend server's keep-alive timeout.

Example

```
ProxyPass "/example" "http://backend.example.com" max=20 ttl=120 retry=300
```

Worker|BalancerMember parameters

Parameter	Default	Description
min	0	Minimum number of connection pool entries, unrelated to the actual number of connections. This only needs to be modified from the default for special circumstances where heap memory associated with the backend connections should be preallocated or retained.
max	1...n	Maximum number of connections that will be allowed to the backend server. The default for this limit is the number of threads per process in the active MPM. In the Prefork MPM, this is always 1, while with other MPMs, it is controlled by the ThreadsPerChild directive.
smax	max	Retained connection pool entries above this limit are freed during certain operations if they have been unused for longer than the time to live, controlled by the ttl parameter. If the connection pool entry has an associated connection, it will be closed. This only needs to be modified from the default for special circumstances where connection pool entries and any associated connections which have exceeded the time to live need to be freed or closed more aggressively.
acquire	-	If set, this will be the maximum time to wait for a free connection in the connection pool, in milliseconds. If there are no free connections in the pool, the Apache httpd will return SERVER_BUSY status to the client.
connectiontimeout	timeout	Connect timeout in seconds. The number of seconds Apache httpd waits for the creation of a connection to the backend to complete. By adding a postfix of ms, the timeout can be also set in milliseconds.
disablereuse	Off	This parameter should be used when you want to force mod_proxy to immediately close a connection to the backend after being used, and thus, disable its persistent connection and pool for that backend. This helps in various situations where a firewall between Apache httpd and the backend server (regardless of protocol) tends to silently drop connections or when backends themselves may be under round- robin DNS. When connection reuse is enabled each backend domain is resolved (with a DNS query) only once per child process and cached for all further connections until the child is recycled. To disable connection reuse, set this property value to On .
enablereuse	On	This is the inverse of 'disablereuse' above, provided as a convenience for scheme handlers that require opt-in for connection reuse (such as mod_proxy_fcgi). 2.4.11 and later only.
flushpackets	off	Determines whether the proxy module will auto-flush the output brigade after each "chunk" of data. 'off' means that it will flush only when needed; 'on' means after each chunk is sent; and 'auto' means poll/wait for a period of time and flush if no input has been received for 'flushwait' milliseconds. Currently, this is in effect only for mod_proxy_ajp and mod_proxy_fcgi.
flushwait	10	The time to wait for additional input, in milliseconds, before flushing the output brigade if 'flushpackets' is 'auto'.
iobuffersize	8192	Adjusts the size of the internal scratchpad IO buffer. This allows you to override the ProxyIOBufferSize for a specific worker. This must be at least 512 or set to 0 for the system default of 8192.
responsefieldsize	8192	Adjust the size of the proxy response field buffer. The buffer size should be at least the size of the largest expected header size from a proxied response. Setting the value to 0 will use the system default of 8192 bytes. Available in Apache HTTP Server 2.4.34 and later.
keepalive	Off	This parameter should be used when you have a firewall between your Apache httpd and the backend server, which tends to drop inactive connections. This flag will tell the Operating System to send KEEP_ALIVE messages on inactive connections and thus prevent the firewall from dropping the connection. To enable keepalive, set this property value to On . The frequency of initial and subsequent TCP keepalive probes depends on global OS settings, and may be as high as 2 hours. To be useful, the frequency configured in the OS must be smaller than the threshold used by the firewall.
lbset	0	Sets the load balancer cluster set that the worker is a member of. The load balancer will try all members of a lower numbered lbset before trying higher numbered ones.
ping	0	Ping property tells the webserver to "test" the connection to the backend before forwarding the request. For AJP, it causes mod_proxy_ajp to send a CPING request on the ajp13 connection (implemented on Tomcat 3.3.2+, 4.1.28+ and 5.0.13+). For HTTP, it causes mod_proxy_http to send a 100-Continue to the backend (only valid for HTTP/1.1 - for non HTTP/1.1 backends, this property has no effect). In both cases, the parameter is the delay in seconds to wait for the reply. This feature has been added to avoid problems with hung and busy backends. This will increase the network traffic during the normal operation which could be an issue, but it will lower the traffic in case some of the cluster nodes are down or busy. By adding a postfix of ms, the delay can be also set in milliseconds.
receivebuffersize	0	Adjusts the size of the explicit (TCP/IP) network buffer size for proxied connections. This allows you to override the ProxyReceiveBufferSize for a specific worker. This must be at least 512 or set to 0 for the system default.
redirect	-	Redirection Route of the worker. This value is usually set dynamically to enable safe removal of the node from the cluster. If set, all requests without session id will be redirected to the BalancerMember that has route parameter equal to this value.
retry	60	Connection pool worker retry timeout in seconds. If the connection pool worker to the backend server is in the error state, Apache httpd will not forward any requests to that server until the timeout expires. This enables to shut down the backend server for maintenance and bring it back online later. A value of 0 means always retry workers in an error state with no timeout.
route	-	Route of the worker when used inside load balancer. The route is a value appended to session id.
status	-	Single letter value defining the initial status of this worker.
		D: Worker is disabled and will not accept any requests.

		S: Worker is administratively stopped. I: Worker is in ignore-errors mode and will always be considered available. R: Worker is a hot spare. For each worker in a given lbset that is unusable (draining, stopped, in error, etc.), a usable hot spare with the same lbset will be used in its place. Hot spares can help ensure that a specific number of workers are always available for use by a balancer. H: Worker is in hot-standby mode and will only be used if no other viable workers or spares are available in the balancer set. E: Worker is in an error state. N: Worker is in drain mode and will only accept existing sticky sessions destined for itself and ignore all other requests.
		Status can be set (which is the default) by prepending with '+' or cleared by prepending with '-'. Thus, a setting of 'S-E' sets this worker to Stopped and clears the in-error flag.
timeout	ProxyTimeout	Connection timeout in seconds. The number of seconds Apache httpd waits for data sent by / to the backend.
ttl	-	Time to live for inactive connections and associated connection pool entries, in seconds. Once reaching this limit, a connection will not be used again; it will be closed at some later time.
flusher	flush	Name of the provider used by <code>mod_proxy_fdpass</code> . See the documentation of this module for more details.
secret	-	Value of secret used by <code>mod_proxy_ajp</code> . It must be identical to the secret configured on the server side of the AJP connection. Available in Apache HTTP Server 2.4.42 and later.
upgrade	-	Protocol accepted by <code>mod_proxy_http</code> or <code>mod_proxy_wstunnel</code> for the HTTP Upgrade mechanism upon negotiation by the HTTP client/browser (per RFC 9110 - Upgrade (↗ https://www.ietf.org/rfc/rfc9110.html#name-upgrade)). See the Protocol Upgrade (↗ #protoupgrade) note below
mapping	-	Type of mapping between the <i>path</i> and the <i>url</i> . This determines the normalization and/or (non-)decoding that <code>mod_proxy</code> will apply to the requested <i>uri-path</i> before matching the <i>path</i> . If a mapping matches, it's committed to the <i>uri-path</i> such that all the directory contexts that use a path (like <code><Location></code>) will be matched using the same mapping. mapping=encoded prevents the %-decoding of the <i>uri-path</i> so that one can use for instance configurations like: <div>ProxyPass "/special%3Fsegment" "https://example.com/special%3Fsegment" mapping=encoded</div> <div><Location "/special%3Fsegment"> Require ip 172.17.2.0/24 </Location></div> mapping=servlet refers to the normalization defined by the Servlet specification, which is for instance applied by Apache Tomcat for servlet containers (notably the path parameters are ignored for the mapping). An <i>uri-path</i> like <code>/some;foo/path</code> is then mapped as <code>/some/path</code> hence matches any of the below regardless of the requested path parameters: <div>ProxyPass "/some/path" "https://servlet.example.com/some/path" mapping=servlet</div> <div><Location "/some/path"> Require valid-user </Location></div> <div>Note It is recommended to use the same mapping on the Apache httpd side than the one used on the backend side. For instance when configuring authorizations in <code><Location></code> blocks for paths that are mapped by <code>mod_proxy</code> to some servlet containers (like applications running on Apache Tomcat), one should use the <code>mapping=servlet</code> setting to prevent path parameters and alike from interfering with the authorizations that are to be enforced in by the Apache httpd.</div>

If the Proxy directive scheme starts with the `balancer://` (eg: `balancer://cluster`, any path information is ignored), then a virtual worker that does not really communicate with the backend server will be created. Instead, it is responsible for the management of several "real" workers. In that case, the special set of parameters can be added to this virtual worker. See `mod_proxy_balancer` for more information about how the balancer works.

Balancer parameters

Parameter	Default	Description
lbmethod	byrequests	Balancer load-balance method. Select the load-balancing scheduler method to use. Either <code>byrequests</code> , to perform weighted request counting; <code>bytraffic</code> , to perform weighted traffic byte count balancing; or <code>bybusyness</code> , to perform pending request balancing. The default is <code>byrequests</code> .
maxattempts	One less than the number of workers, or 1 with a single worker.	Maximum number of failover attempts before giving up.
nofailover	Off	If set to <code>On</code> , the session will break if the worker is in error state or disabled. Set this value to <code>On</code> if backend servers do not support session replication.
stickysession	-	Balancer sticky session name. The value is usually set to something like <code>JSESSIONID</code> or <code>PHPSESSIONID</code> , and it depends on the backend application server that support sessions. If the backend application server uses different name for cookies and url encoded id (like servlet containers) use <code> </code> to separate them. The first part is for the cookie the second for the path. Available in Apache HTTP Server 2.4.4 and later.
stickysessionsep	"."	Sets the separation symbol in the session cookie. Some backend application servers do not use the <code>'.'</code> as the symbol. For example, the Oracle Weblogic server uses <code>'!'</code> . The correct symbol can be set using this option. The setting of 'Off' signifies that no symbol is used.
scolonpathdelim	Off	If set to <code>On</code> , the semi-colon character <code>';</code> will be used as an additional sticky session path delimiter/separator. This is mainly used to emulate <code>mod_jk</code> 's behavior when dealing with paths such as <code>JSESSIONID=6736bcf34;foo=aabfa</code>
timeout	0	Balancer timeout in seconds. If set, this will be the maximum time to wait for a free worker. The default is to not wait.
failonstatus	-	A single or comma-separated list of HTTP status codes. If set, this will force the worker into error state when the backend returns any status code in the list. Worker recovery behaves the same as other worker errors.
failontimeout	Off	If set, an IO read timeout after a request is sent to the backend will force the worker into error state. Worker recovery behaves the same as other worker errors. Available in Apache HTTP Server 2.4.5 and later.
nonce	<auto>	The protective nonce used in the <code>balancer-manager</code> application page. The default is to use an automatically determined UUID-based nonce, to provide for further protection for the page. If set, then the nonce is set to that value. A setting of <code>NONE</code> disables all nonce checking.

		<div>Note</div> <div>In addition to the nonce, the balancer -manager page should be protected via an ACL.</div>
growth	0	Number of additional BalancerMembers to allow to be added to this balancer in addition to those defined at configuration.
forcerecovery	On	Force the immediate recovery of all workers without considering the retry parameter of the workers if all workers of a balancer are in error state. There might be cases where an already overloaded backend can get into deeper trouble if the recovery of all workers is enforced without considering the retry parameter of each worker. In this case, set to Off. Available in Apache HTTP Server 2.4.2 and later.

A sample balancer setup:

```
ProxyPass "/special-area" "http://special.example.com" smax=5 max=10
ProxyPass "/" "balancer://mycluster/" stickysession=JSESSIONID|jsessionid nofailover=On
<Proxy "balancer://mycluster">
  BalancerMember "ajp://1.2.3.4:8009"
  BalancerMember "ajp://1.2.3.5:8009" loadfactor=20
  # Less powerful server, don't send as many requests there,
  BalancerMember "ajp://1.2.3.6:8009" loadfactor=5
</Proxy>
```

Configuring hot spares can help ensure that a certain number of workers are always available for use per load balancer set:

```
ProxyPass "/" "balancer://sparecluster/"
<Proxy balancer://sparecluster>
  BalancerMember ajp://1.2.3.4:8009
  BalancerMember ajp://1.2.3.5:8009
  # The servers below are hot spares. For each server above that is unusable
  # (draining, stopped, unreachable, in error state, etc.), one of these spares
  # will be used in its place. Two servers will always be available for a request
  # unless one or more of the spares is also unusable.
  BalancerMember ajp://1.2.3.6:8009 status=+R
  BalancerMember ajp://1.2.3.7:8009 status=+R
</Proxy>
```

Setting up a hot-standby that will only be used if no other members (or spares) are available in the load balancer set:

```
ProxyPass "/" "balancer://hotcluster/"
<Proxy "balancer://hotcluster">
  BalancerMember "ajp://1.2.3.4:8009" loadfactor=1
  BalancerMember "ajp://1.2.3.5:8009" loadfactor=2.25
  # The server below is on hot standby
  BalancerMember "ajp://1.2.3.6:8009" status=+H
  ProxySet lbmethod=bytraffic
</Proxy>
```

Additional ProxyPass Keywords

Normally, mod_proxy will canonicalise ProxyPassed URLs. But this may be incompatible with some backends, particularly those that make use of *PATH_INFO*. The optional *nocanon* keyword suppresses this and passes the URL path "raw" to the backend. Note that this keyword may affect the security of your backend, as it removes the normal limited protection against URL-based attacks provided by the proxy.

Normally, mod_proxy will include the query string when generating the *SCRIPT_FILENAME* environment variable. The optional *noquery* keyword (available in httpd 2.4.1 and later) prevents this.

The optional *interpolate* keyword, in combination with **ProxyPassInterpolateEnv**, causes the ProxyPass to interpolate environment variables, using the syntax *\${VARNAME}*. Note that many of the standard CGI-derived environment variables will not exist when this interpolation happens, so you may still have to resort to **mod_rewrite** for complex rules. Also note that interpolation is supported within the scheme/hostname/port portion of a URL only for variables that are available when the directive is parsed (like **Define**). Dynamic determination of those fields can be accomplished with **mod_rewrite**. The following example describes how to use **mod_rewrite** to dynamically set the scheme to http or https:

```
RewriteEngine On

RewriteCond "%{HTTPS}" =off
RewriteRule "." "-" [E=protocol:http]
RewriteCond "%{HTTPS}" =on
RewriteRule "." "-" [E=protocol:https]

RewriteRule "^/mirror/foo/(.*)" "%{ENV:protocol}://backend.example.com/$1" [P]
ProxyPassReverse "/mirror/foo/" "http://backend.example.com/"
ProxyPassReverse "/mirror/foo/" "https://backend.example.com/"
```

Protocol Upgrade

Since Apache HTTP Server 2.4.47, protocol Upgrade (tunneling) can be handled end-to-end by **mod_proxy_http** using the **ProxyPass** parameter *upgrade*. End-to-end means that the HTTP Upgrade request from the client/browser is first forwarded by **mod_proxy_http** to the origin server and the connection will be upgraded (and tunneled by **mod_proxy_http**) only if the origin server accepts/initiates the upgrade (HTTP response 101 **Switching Protocols**). If the origin server responds with anything else **mod_proxy_http** will continue forwarding (and enforcing) the HTTP protocol as usual for this connection. See Websocket Upgrade (2.4.47 and later) ([↗ #wsupgrade](#)) for an example of configuration using **mod_proxy_http**. For Apache HTTP Server 2.4.46 and earlier (or if **ProxyWebsocketFallbackToProxyHttp** from 2.4.48 and later disables **mod_proxy_http** handling), see the documentation of **mod_proxy_wstunnel** for how to proxy the WebSocket protocol.

ProxyPassInherit Directive

Description:	Inherit ProxyPass directives defined from the main server
Syntax:	ProxyPassInherit On Off
Default:	ProxyPassInherit On
Context:	server config, virtual host

Status:	Extension
Module:	mod_proxy
Compatibility:	ProxyPassInherit is only available in Apache HTTP Server 2.4.5 and later.

This directive will cause the current server/vhost to "inherit" **ProxyPass** directives defined in the main server. This can cause issues and inconsistent behavior if using the Balancer Manager for dynamic changes and so should be disabled if using that feature.

The setting in the global server defines the default for all vhosts.

Disabling ProxyPassInherit also disables **BalancerInherit**.

ProxyPassInterpolateEnv Directive

Description:	Enable Environment Variable interpolation in Reverse Proxy configurations
Syntax:	ProxyPassInterpolateEnv On Off
Default:	ProxyPassInterpolateEnv Off
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy
Compatibility:	Available in httpd 2.2.9 and later

This directive, together with the `interpolate` argument to **ProxyPass**, **ProxyPassReverse**, **ProxyPassReverseCookieDomain**, and **ProxyPassReverseCookiePath**, enables reverse proxies to be dynamically configured using environment variables which may be set by another module such as **mod_rewrite**. It affects the **ProxyPass**, **ProxyPassReverse**, **ProxyPassReverseCookieDomain**, and **ProxyPassReverseCookiePath** directives and causes them to substitute the value of an environment variable `varname` for the string `${varname}` in configuration directives if the `interpolate` option is set.

The scheme/hostname/port portion of **ProxyPass** may contain variables, but only the ones available when the directive is parsed (for example, using **Define**). For all the other use cases, please consider using **mod_rewrite** instead.

Performance warning

Keep this turned off unless you need it! Adding variables to **ProxyPass** for example may lead to the use of the default mod_proxy's workers configured (that don't allow any fine tuning like connections reuse, etc..).

ProxyPassMatch Directive

Description:	Maps remote servers into the local server URL-space using regular expressions
Syntax:	ProxyPassMatch [<i>regex</i>] <code>! url</code> [<i>key=value</i> [<i>key=value ...</i>]]
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy

This directive is equivalent to **ProxyPass** but makes use of regular expressions instead of simple prefix matching. The supplied regular expression is matched against the *url*, and if it matches, the server will substitute any parenthesized matches into the given string and use it as a new *url*.

Note: This directive cannot be used within a `<Directory>` context.

Suppose the local server has address `http://example.com/`; then

```
ProxyPassMatch "^/(.*\.gif)$" "http://backend.example.com/$1"
```

will cause a local request for `http://example.com/foo/bar.gif` to be internally converted into a proxy request to `http://backend.example.com/foo/bar.gif`.

Note

The URL argument must be parsable as a URL *before* regexp substitutions (as well as after). This limits the matches you can use. For instance, if we had used

```
ProxyPassMatch "^/(.*\.gif)$" "http://backend.example.com:8000$1"
```

in our previous example, it would fail with a syntax error at server startup. This is a bug (PR 46665 in the ASF bugzilla), and the workaround is to reformulate the match:

```
ProxyPassMatch "^/(.*\.gif)$" "http://backend.example.com:8000/$1"
```

The `!` directive is useful in situations where you don't want to reverse-proxy a subdirectory.

When used inside a `<LocationMatch>` section, the first argument is omitted and the regexp is obtained from the `<LocationMatch>`.

If you require a more flexible reverse-proxy configuration, see the **RewriteRule** directive with the `[P]` flag.

Default Substitution

When the URL parameter doesn't use any backreferences into the regular expression, the original URL will be appended to the URL parameter.

Security Warning

Take care when constructing the target URL of the rule, considering the security impact from allowing the client influence over the set of URLs to which your server will act as a proxy. Ensure that the scheme and hostname part of the URL is either fixed or does not allow the client undue influence.

ProxyPassReverse Directive

Description:	Adjusts the URL in HTTP response headers sent from a reverse proxied server
Syntax:	ProxyPassReverse [<i>path</i>] <i>url</i> [<i>interpolate</i>]
Context:	server config, virtual host, directory

Status:	Extension
Module:	mod_proxy

This directive lets Apache httpd adjust the URL in the `Location`, `Content-Location` and `URI` headers on HTTP redirect responses. This is essential when Apache httpd is used as a reverse proxy (or gateway) to avoid bypassing the reverse proxy because of HTTP redirects on the backend servers which stay behind the reverse proxy.

Only the HTTP response headers specifically mentioned above will be rewritten. Apache httpd will not rewrite other response headers, nor will it by default rewrite URL references inside HTML pages. This means that if the proxied content contains absolute URL references, they will bypass the proxy. To rewrite HTML content to match the proxy, you must load and enable `mod_proxy_html`.

path is the name of a local virtual path; *url* is a partial URL for the remote server. These parameters are used the same way as for the `ProxyPass` directive.

For example, suppose the local server has address `http://example.com/`; then

```
ProxyPass      "/mirror/foo/" "http://backend.example.com/"
ProxyPassReverse "/mirror/foo/" "http://backend.example.com/"
ProxyPassReverseCookieDomain "backend.example.com" "public.example.com"
ProxyPassReverseCookiePath  "/" "/mirror/foo/"
```

will not only cause a local request for the `http://example.com/mirror/foo/bar` to be internally converted into a proxy request to `http://backend.example.com/bar` (the functionality which `ProxyPass` provides here). It also takes care of redirects which the server `backend.example.com` sends when redirecting `http://backend.example.com/bar` to `http://backend.example.com/quux`. Apache httpd adjusts this to `http://example.com/mirror/foo/quux` before forwarding the HTTP redirect response to the client. Note that the hostname used for constructing the URL is chosen in respect to the setting of the `UseCanonicalName` directive.

Note that this `ProxyPassReverse` directive can also be used in conjunction with the proxy feature (`RewriteRule ... [P]`) from `mod_rewrite` because it doesn't depend on a corresponding `ProxyPass` directive.

The optional `interpolate` keyword, used together with `ProxyPassInterpolateEnv`, enables interpolation of environment variables specified using the format `${VARNAME}`. Note that interpolation is not supported within the scheme portion of a URL.

When used inside a `<Location>` section, the first argument is omitted and the local directory is obtained from the `<Location>`. The same occurs inside a `<LocationMatch>` section, but will probably not work as intended, as `ProxyPassReverse` will interpret the regexp literally as a path; if needed in this situation, specify the `ProxyPassReverse` outside the section or in a separate `<Location>` section.

This directive is not supported in `<Directory>` or `<Files>` sections.

ProxyPassReverseCookieDomain Directive

Description:	Adjusts the Domain string in Set-Cookie headers from a reverse- proxied server
Syntax:	<code>ProxyPassReverseCookieDomain <i>internal-domain public-domain</i> [interpolate]</code>
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy

Usage is basically similar to `ProxyPassReverse`, but instead of rewriting headers that are a URL, this rewrites the domain string in `Set-Cookie` headers.

ProxyPassReverseCookiePath Directive

Description:	Adjusts the Path string in Set-Cookie headers from a reverse- proxied server
Syntax:	<code>ProxyPassReverseCookiePath <i>internal-path public-path</i> [interpolate]</code>
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy

Useful in conjunction with `ProxyPassReverse` in situations where backend URL paths are mapped to public paths on the reverse proxy. This directive rewrites the `path` string in `Set-Cookie` headers. If the beginning of the cookie path matches *internal-path*, the cookie path will be replaced with *public-path*.

In the example given with `ProxyPassReverse`, the directive:

```
ProxyPassReverseCookiePath  "/" "/mirror/foo/"
```

will rewrite a cookie with backend path `/` (or `/example` or, in fact, anything) to `/mirror/foo/`.

ProxyPreserveHost Directive

Description:	Use incoming Host HTTP request header for proxy request
Syntax:	<code>ProxyPreserveHost On Off</code>
Default:	<code>ProxyPreserveHost Off</code>
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy
Compatibility:	Usable in directory context in 2.3.3 and later.

When enabled, this option will pass the `Host :` line from the incoming request to the proxied host, instead of the hostname specified in the `ProxyPass` line.

This option should normally be turned `Off`. It is mostly useful in special configurations like proxied mass name-based virtual hosting, where the original `Host` header needs to be evaluated by the backend server.

ProxyReceiveBufferSize Directive

Description:	Network buffer size for proxied HTTP and FTP connections
Syntax:	<code>ProxyReceiveBufferSize <i>bytes</i></code>
Default:	<code>ProxyReceiveBufferSize 0</code>
Context:	server config, virtual host

Status:	Extension
Module:	mod_proxy

The **ProxyReceiveBufferSize** directive specifies an explicit (TCP/IP) network buffer size for proxied HTTP and FTP connections, for increased throughput. It has to be greater than 512 or set to 0 to indicate that the system's default buffer size should be used.

Example

ProxyReceiveBufferSize 2048

ProxyRemote Directive

Description:	Remote proxy used to handle certain requests
Syntax:	ProxyRemote <i>match remote-server</i>
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

This defines remote proxies to this proxy. *match* is either the name of a URL-scheme that the remote server supports, or a partial URL for which the remote server should be used, or * to indicate the server should be contacted for all requests. *remote-server* is a partial URL for the remote server. Syntax:

```
remote-server = scheme://hostname[:port]
```

scheme is effectively the protocol that should be used to communicate with the remote server; only http and https are supported by this module. When using https, the requests are forwarded through the remote proxy using the HTTP CONNECT method.

Example

ProxyRemote "http://goodguys.example.com/" "http://mirrorguys.example.com:8000"
ProxyRemote "*" "http://cleverproxy.localdomain"
ProxyRemote "ftp" "http://ftpproxy.mydomain:8080"

In the last example, the proxy will forward FTP requests, encapsulated as yet another HTTP proxy request, to another proxy which can handle them.

This option also supports reverse proxy configuration; a backend webserver can be embedded within a virtualhost URL space even if that server is hidden by another forward proxy.

ProxyRemoteMatch Directive

Description:	Remote proxy used to handle requests matched by regular expressions
Syntax:	ProxyRemoteMatch <i>regex remote-server</i>
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

The **ProxyRemoteMatch** is identical to the **ProxyRemote** directive, except that the first argument is a [regular expression \(↗ ../glossary.html#regex\)](#) match against the requested URL.

ProxyRequests Directive

Description:	Enables forward (standard) proxy requests
Syntax:	ProxyRequests On Off
Default:	ProxyRequests Off
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

This allows or prevents Apache httpd from functioning as a forward proxy server. (Setting ProxyRequests to Off does not disable use of the **ProxyPass** directive.)

In a typical reverse proxy or gateway configuration, this option should be set to Off.

In order to get the functionality of proxying HTTP or FTP sites, you need also **mod_proxy_http** or **mod_proxy_ftp** (or both) present in the server.

In order to get the functionality of (forward) proxying HTTPS sites, you need **mod_proxy_connect** enabled in the server.

Warning

Do not enable proxying with **ProxyRequests** until you have secured your server (↗ #access) . Open proxy servers are dangerous both to your network and to the Internet at large.

See also

- Forward and Reverse Proxies/Gateways

ProxySet Directive

Description:	Set various Proxy balancer or member parameters
Syntax:	ProxySet <i>url key=value [key=value ...]</i>
Context:	server config, virtual host, directory
Status:	Extension
Module:	mod_proxy
Compatibility:	ProxySet is only available in Apache HTTP Server 2.2 and later.

This directive is used as an alternate method of setting any of the parameters available to Proxy balancers and workers normally done via the **ProxyPass** directive. If used within a <Proxy *balancer url|worker url*> container directive, the *url* argument is not required. As a side effect the respective balancer or worker gets created. This can be useful when doing reverse proxying via a **RewriteRule** instead of a **ProxyPass** directive.

```
<Proxy "balancer://hotcluster">
  BalancerMember "http://www2.example.com:8080" loadfactor=1
  BalancerMember "http://www3.example.com:8080" loadfactor=2
  ProxySet lbmethod=bytraffic
</Proxy>
```

```
<Proxy "http://backend">
  ProxySet keepalive=On
</Proxy>
```

```
ProxySet "balancer://foo" lbmethod=bytraffic timeout=15
```

```
ProxySet "ajp://backend:7001" timeout=15
```

Warning

Keep in mind that the same parameter key can have a different meaning depending whether it is applied to a balancer or a worker, as shown by the two examples above regarding timeout.

ProxySourceAddress Directive

Description:	Set local IP address for outgoing proxy connections
Syntax:	ProxySourceAddress <i>address</i>
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy
Compatibility:	Available in version 2.3.9 and later

This directive allows to set a specific local address to bind to when connecting to a backend server.

ProxyStatus Directive

Description:	Show Proxy LoadBalancer status in mod_status
Syntax:	ProxyStatus Off On Full
Default:	ProxyStatus Off
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy
Compatibility:	Available in version 2.2 and later

This directive determines whether or not proxy loadbalancer status data is displayed via the `mod_status` server-status page.

Note

Full is synonymous with On

ProxyTimeout Directive

Description:	Network timeout for proxied requests
Syntax:	ProxyTimeout <i>seconds</i>
Default:	Value of Timeout
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

This directive allows a user to specify a timeout on proxy requests. This is useful when you have a slow/buggy appserver which hangs, and you would rather just return a timeout and fail gracefully instead of waiting however long it takes the server to return.

ProxyVia Directive

Description:	Information provided in the Via HTTP response header for proxied requests
Syntax:	ProxyVia On Off Full Block
Default:	ProxyVia Off
Context:	server config, virtual host
Status:	Extension
Module:	mod_proxy

This directive controls the use of the `Via :` HTTP header by the proxy. Its intended use is to control the flow of proxy requests along a chain of proxy servers. See RFC 2616 ([↗ http://www.ietf.org/rfc/rfc2616.txt](http://www.ietf.org/rfc/rfc2616.txt)) (HTTP/1.1), section 14.45 for an explanation of `Via :` header lines.

- If set to `Off`, which is the default, no special processing is performed. If a request or reply contains a `Via :` header, it is passed through unchanged.
- If set to `On`, each request and reply will get a `Via :` header line added for the current host.
- If set to `Full`, each generated `Via :` header line will additionally have the Apache httpd server version shown as a `Via :` comment field.
- If set to `Block`, every proxy request will have all its `Via :` header lines removed. No new `Via :` header will be generated.

Comments

Notice:

This is not a Q&A section. Comments placed here should be pointed towards suggestions on improving the documentation or server, and may be removed by our moderators if they are either implemented or considered invalid/off-topic. Questions on how to manage the Apache HTTP Server should be directed at either our IRC channel, #httpd, on Libera.chat, or sent to our mailing lists.

